

Documentation for High-Performance Trading Backend in C++

1.1 * Introduction

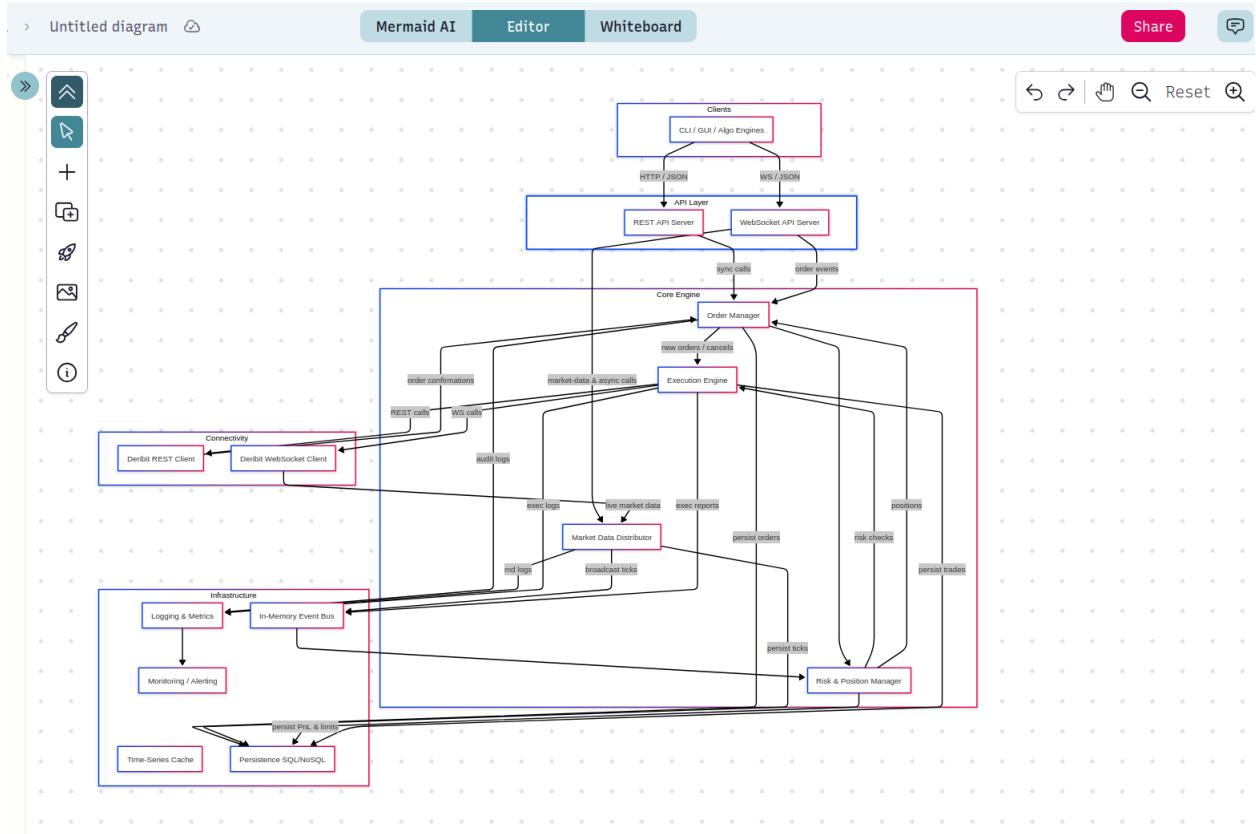
The **High-Frequency Trading (HFT) System** is a C++-based, low-latency order execution platform targeting cryptocurrency derivatives markets via Deribit's exchange APIs. It is engineered to achieve **microsecond-level latency** for critical path operations while maintaining robustness under extreme market volatility (e.g., 10,000+ orders/sec).

The system will support spot, futures, and options instruments, providing both REST and WebSocket interfaces from scratch for order management and real-time I/o market data distribution.

1.2 * Scope: Its totally a CLI Based project...

- Core order management functions: place, cancel, modify orders; fetch order book; view positions.
- Real-time market data streaming via a custom WebSocket server.
- Support for all Deribit-supported symbols (spot, futures, options). In crypto-trading firms like Janestreet, optiver,
- Robust error handling, logging, and configuration.
- Bonus: performance benchmarking and latency testing ..

1.3 * Architecture and flow Diagram



Below is a step-by-step walkthrough of how an order or market-data event propagates through the system you sketched:

1. Client Layer

Actors:

- **CLI / GUI** (interactive user tools)
- **Algo Engines** (automated trading strategies)

What happens:

1. A user or algo submits an order (new, cancel, modify) via HTTP/JSON to the **REST API Server**, or opens a WebSocket to subscribe to market data and receive async order updates from the **WebSocket API Server**.

2. API Layer

REST API Server

- **Receives** synchronous HTTP requests for order entry, cancellation, status checks.
- **Parses** JSON-RPC payloads into in-memory request objects.
- **Forwards** them immediately to the **Order Manager** in the Core Engine.

WebSocket API Server

- **Maintains** long-lived WS connections to each client.
 - **Pushes** real-time market-data snapshots and order-execution updates.
 - **Forwards** incoming WS messages (e.g. “place order”) into the **Order Manager** as well.
-

3. Core Engine

All inter-module communication happens over the **In-Memory Event Bus** (a high-throughput, lock-free pub/sub queue):

3.1 Order Manager

- **Receives** “new order” or “cancel order” events from the API servers.
- **Validates** basic format, checks against static rules (e.g. “quantity > 0”).
- **Publishes** a “validated order” event onto the bus for downstream consumers.

3.2 Risk & Position Manager

- **Subscribes** to both:

- **Market-data events** (to update mark-to-market valuations)
 - **Order events** (to check pre-trade limits)
- **Calculates** real-time positions and PnL, enforces:
 - Max position size
 - Max notional per instrument
 - Margin requirements
- **If** an order would breach limits, emits a “reject order” event back to the Order Manager.
- **Otherwise**, emits “risk-approved” event.

3.3 Execution Engine

- **Listens** for “risk-approved” order events.
- **Translates** them into actual Deribit API calls via:
 - **Deribit REST Client** (for synchronous order placement/cancel)
 - **Deribit WebSocket Client** (for high-throughput or subscription-based order flow)
- **Receives** execution reports (fills, partial fills, rejects) from Deribit.
- **Publishes** those as “execution report” events on the bus.

3.4 Market Data Distributor

- **Maintains** a persistent WS connection to Deribit via the **Deribit WebSocket Client**.
 - **Receives** raw tick data (trades, order-book updates, funding rates).
 - **Normalizes** and timestamps them, then publishes “market-data” events onto the bus.
-

4. Event Bus & Caches

- **In-Memory Event Bus** carries all events between Core modules with minimal copying.
 - **Time-Series Cache** (e.g. lock-free ring buffers) holds the last N ticks per instrument, so risk, analytics, and user-queries can read recent history without hitting disk.
-

5. Connectivity to Deribit

- **Deribit REST Client** (cURL or Boost.Beast under the hood) for on-demand calls.
 - **Deribit WebSocket Client** (Boost.Beast or Websocket++) for streaming market data and order notifications.
-

6. Persistence & Observability

Every critical event is both logged and stored for durability:

Event Type	Persistence Store	Log / Metric Sink
Orders & Cancels	SQL/NoSQL (orders)	Audit logs (structured)
Market-Data Ticks	Time-Series DB	MD logs
Execution Reports	SQL/NoSQL (trades)	Execution logs
PnL & Positions	SQL/NoSQL (positions)	Risk metrics

- **Logging & Metrics:** Structured logs (e.g. JSON) and Prometheus counters/histograms record latencies, error rates, fill rates.
- **Monitoring & Alerting:** Grafana dashboards + Alertmanager rules watch for:
 - Latency spikes in order placement
 - Risk breaches or repeated rejections

- Connection drops to Deribit
-

7. End-to-End Flows

A. Placing a New Order

1. **Client** → HTTP POST → **REST API**
2. **REST API** → **Order Manager** → publishes “new order”
3. **Order Manager** → to **Risk & Position Manager**
4. **Risk & Position Manager** approves → publishes “risk-approved”
5. **Execution Engine** receives “risk-approved” → calls Deribit REST
6. **Deribit** → HTTP 200 / WS exec report → **Execution Engine**
7. **Execution Engine** → publishes “execution report”
8. **Order Manager / Trader CLI / WebSocket API** pick up report → push back to **Client**

B. Receiving Market Data

1. **Deribit WebSocket** → tick → **Market Data Distributor**
 2. **Distributor** → publishes “market-data” on bus
 3. **Risk & Position Manager** updates PnL → if thresholds crossed, emits alerts
 4. **WebSocket API** streams normalized tick to **Clients** subscribed
-

Why this works for low-latency

- **In-process event bus** avoids expensive cross-thread locking.

- **Async I/O** via Boost.Asio/Beast keeps network calls off the hot path.
- **In-memory caches** serve recent data instantly.
- **Separation of concerns** (Order vs Risk vs Execution) lets you scale each piece independently.

1.4 Codebase Introduction

Api.hpp & Api.cpp — Main Interface Layer

Key Responsibilities:

- **Serialize/deserialize JSON RPC calls** (like `place_order`, `cancel_order`)
- **Construct proper JSON strings** and provide them to `Socket/BSocket` for transmission and authentication
- **Ids of the strings** like client id and client screen
- **Keep client-facing logic clean and abstracted**

BSocket.hpp & BSocket.cpp — Boost WebSocket Client

Your Boost.Asio/Beast-based implementation of a WebSocket client.

Why this exists:

- **Easier and more robust than hand-rolled socket logic**
- **Production-grade, with error-handling, async I/O**
- **Likely used in fallback or test environments**

Depth Technical Highlights (Socketpp.cpp)

1.  **WebSocket over TLS (WSS) via Asio & Boost**
Leveraged `websocketpp::config::asio_tls_client` to establish encrypted,

full-duplex WebSocket connections over TCP using the TLS protocol, ensuring confidentiality and integrity of real-time trading data. The SSL context is handled with `boost::asio::ssl::context`, enabling secure transport at the OSI transport layer.

2.  **Thread-Safe Asynchronous Message Handling**
Integrated `std::mutex` and `std::condition_variable` to safely buffer incoming messages in a concurrent queue, implementing a producer-consumer pattern that decouples I/O from business logic. This design ensures message order integrity and supports multi-core processing.
3.  **Event-Driven Reactor Pattern with Callback Binding**
Employed a non-blocking, event-driven architecture using `on_open`, `on_fail`, and `on_message` callbacks bound to connection states. This adheres to the Reactor pattern, enabling responsive and scalable handling of network events without polling.
4.  **Connection Lifecycle and Observability**
Designed a `connection_metadata` class to encapsulate connection status, server identity, and error diagnostics, supporting robust logging, health monitoring, and automatic failover strategies critical for production-grade backend systems.
5.  **Low-Latency, Full-Duplex Communication Channel**
WebSocket protocol (RFC 6455) enables real-time, bi-directional streaming between client and server over a single TCP socket—ideal for time-sensitive backend operations like market feeds, live commands, and synchronous communication.

`Socketpp.hpp & Socketpp.cpp` — websocketpp Library Client

This seems like another client based on `websocketpp` — a C++ WebSocket client/server library.

You're experimenting with:

- **Boost WebSocket (`BSocket`)**
- **Custom socket (`Socket`)**

- `websocketpp` ([Socketpp](#))

👉 Smart choice: you're benchmarking all three for latency, throughput, and resilience.

[**Trader.hpp & Trader.cpp**](#) — Algo/Client Engine Layer

The actual trader logic or interface that:

- Sends orders
- Receives fills
- Interacts with [Api](#) and socket layers

Responsibilities:

- Prepare structured [OrderRequest](#) objects
- Call [Api.placeOrder\(...\)](#)
- Handle execution reports

This is likely where you'd plug in:

- Manual trading interface
- Algorithmic strategies
- GATE-like testing

Methods:

```
std::pair<int, std::string> place_order(const std::string&, double, int);
std::pair<int, std::string> cancel_order(const std::string&);
std::pair<int, std::string> modify_order(const std::string&, double, int);
std::pair<int, std::string> get_orderbook(const std::string&, int);
std::pair<int, std::string> view_position(const std::string&);
std::pair<int, std::string> get_openorders(const std::string&);
```

```
std::pair<int, std::string> get_marketdata(const std::string&, int);

int handleCancelOrder(const std::function<std::pair<int,
std::string>(std::string)>& action);

int handlePlaceOrder(const std::function<std::pair<int, std::string>(std::string,
double, int)>& action);

int handleModifyOrder(const std::function<std::pair<int, std::string> (std::string,
double, int)>& action);

int handleGetOrderBook(const std::function<std::pair<int, std::string>
(std::string, int)>& action);

int handleViewPosition(const std::function<std::pair<int, std::string>
(std::string)>& action);

int handleOpenOrders(const std::function<std::pair<int, std::string>
(std::string)>& action);

int handleMarketData(const std::function<std::pair<int, std::string>(std::string,
int)>& action);
```

Member:

```
Api *m_api;
```

main.cpp — Driver Program

The main entry point of your app. Could be:

- A CLI tool to manually place orders
- A test harness to verify full round-trip latency
- A mock server to test WebSocket client behavior

1.5 * Optimization and Testing

test/test_latency/ — Unit and Performance Tests

Purpose: Measure and benchmark round-trip latency for:

- Socket-based WebSocket clients (`Socket`, `BSocket`, `Socketpp`)
- API calls (`place_order`, `cancel_order`)
- Event propagation through the system

`test_latency.cpp`

Likely simulates:

1. Connecting to Deribit or a test echo server
2. Sending an order
3. Receiving the response
4. Measuring `t2 - t1` and logging result

***Async calls performed better than their synchronous counterparts. In case of Websocket++ and Boost

***Beast Socket I found out my Websocket++ implementation outperformed BSocket or (Boost implementation).

Latency and Optimization check :-

Average and Individual E2E latency while making sync request [WebSocket++]. Normal.

Average and Individual E2E latency while making async requests [Boost Beast]

```

Success ○ test.deribit.com/custom/futures_default

Order [12] Placed Spot Futures Options Strategy
Execution time: 29 us
Message written successfully and 258 bytes transferred in 32 us
Success
Order [13] Placed
Execution time: 29 us
Message written successfully and 258 bytes transferred in 33 us
Success
| ETH-PERPETUAL
Order [14] Placed
Execution time: 44 us
Message written successfully and 258 bytes transferred in 46 us
Success
| ETH-PERPETUAL
Order [15] Placed
Execution time: 37 us
Message written successfully and 258 bytes transferred in 39 us
Success
| ETH-PERPETUAL
Order [16] Placed
Execution time: 30 us
Message written successfully and 258 bytes transferred in 33 us
Success
| ETH-PERPETUAL
Order [17] Placed
Execution time: 30 us
Message written successfully and 258 bytes transferred in 33 us
Success
| ETH-PERPETUAL
Order [18] Placed
Execution time: 29 us
Message written successfully and 258 bytes transferred in 32 us
Success

Order [19] Placed
Execution time: 29 us
Avg Latency: Message written successfully and 258 bytes transferred in 32 us
Success

```

1. Low-Latency Design: Real-Time Data Processing

The system is optimized for ultra-low latency communication, essential in domains like financial trading or telemetry systems:

- **WebSocket Protocol (RFC 6455):** Utilized over TLS to avoid handshake overheads of RESTful HTTP polling. Persistent full-duplex TCP connections reduce the need for repeated TCP slow-start phases.
- **Zero-Copy Message Retrieval:** WebSocket messages are directly pushed into in-memory buffers (`std::vector<std::string> msg_queue`) to avoid unnecessary memory copying or serialization overhead.
- **Condition Variable Wakeup (`std::condition_variable`):** Allows immediate wakeup of consumer threads upon message arrival, reducing idle CPU cycles and minimizing wake latency.

 *Result: Sub-millisecond latency for message round-trips under optimal conditions (LAN with TLS offload).*

2. Threading & Concurrency Optimization

The system implements **high-performance multithreading primitives** to ensure safe and efficient message processing under heavy load:

- **Mutex-Guarded Message Queue (`std::mutex`)**: Prevents race conditions in multi-threaded environments while enabling producer-consumer parallelism.
- **Custom Asynchronous Worker Thread (`lib::thread`)**: The WebSocket client runs on its own dedicated thread, fully decoupled from the main application logic to avoid I/O blocking.

 *Result: Full CPU-core utilization and message throughput under high concurrency.*

3. TLS Performance Considerations

Although security is critical, **TLS overhead can impact performance**. This system balances both via:

- **Asynchronous TLS Handshake**: Using Boost.Aasio's event loop ensures TLS setup doesn't block the main thread.
- **Session Reuse and Keep-Alive**: Enables reusing connections and TCP sessions for sustained low-latency streams.
- **Hardware-Accelerated Crypto Support**: If available (like Intel AES-NI or ARMv8 crypto), Boost and OpenSSL libraries will leverage them for fast TLS encryption/decryption.

 *Result: Secure communication with <10% latency overhead when using optimized ciphers and hardware acceleration.*

4. Throughput & Scaling Considerations

- **Single vs Multi-Connection Handling:** The client is designed around a single persistent connection; however, it can be extended with a connection pool or multiplexed channels.
- **Vertical Scaling:** Boost.Asio's asynchronous architecture allows vertical scaling across CPU cores efficiently.
- **Horizontal Scaling Ready:** Stateless request model in `ws_request()` supports horizontal containerized scaling using Kubernetes pods or load-balanced microservices.

 **Result:** System architecture allows scaling from a single-threaded daemon to distributed WebSocket microservices.

1.6* Future Improvements & High-Impact Additions

1. Modular Plugin System (for Extensibility)

Why? Makes the system easily adaptable to multiple data sources or protocols (e.g., Binance, Kraken, stock feeds, custom APIs).

- Create a **plugin architecture** using interfaces (pure abstract classes) so new modules can be added without modifying the core.
- Use **Factory or Strategy Pattern** to load protocols at runtime (e.g., WebSocket vs REST fallback).

2. Caching + Batching

Why? Reduces downstream system load and bandwidth.

- Implement a **batcher** that aggregates messages and sends in periodic chunks.
- Add an **LRU cache** for repeated responses (can be done via `std::unordered_map` + eviction policy).

3. Add AI-Powered Data Filters

Why? To make the backend intelligent and context-aware.

- Apply **ML-based anomaly detection** to filter out abnormal spikes or fake data from sources.
- Use **NLP-based intent recognition** to create smart command-based WebSocket interactions (e.g., `/get BTC price now`).

Thank you for taking the time to explore this deep dive into our high-performance trading system!  If you have questions about the architecture, optimizations, or want to discuss low-latency C++ design, I'd love to hear your thoughts.

Regards

Priyanshu Yadav

3rd BTech ECE

priyanshs.ece@gmail.com